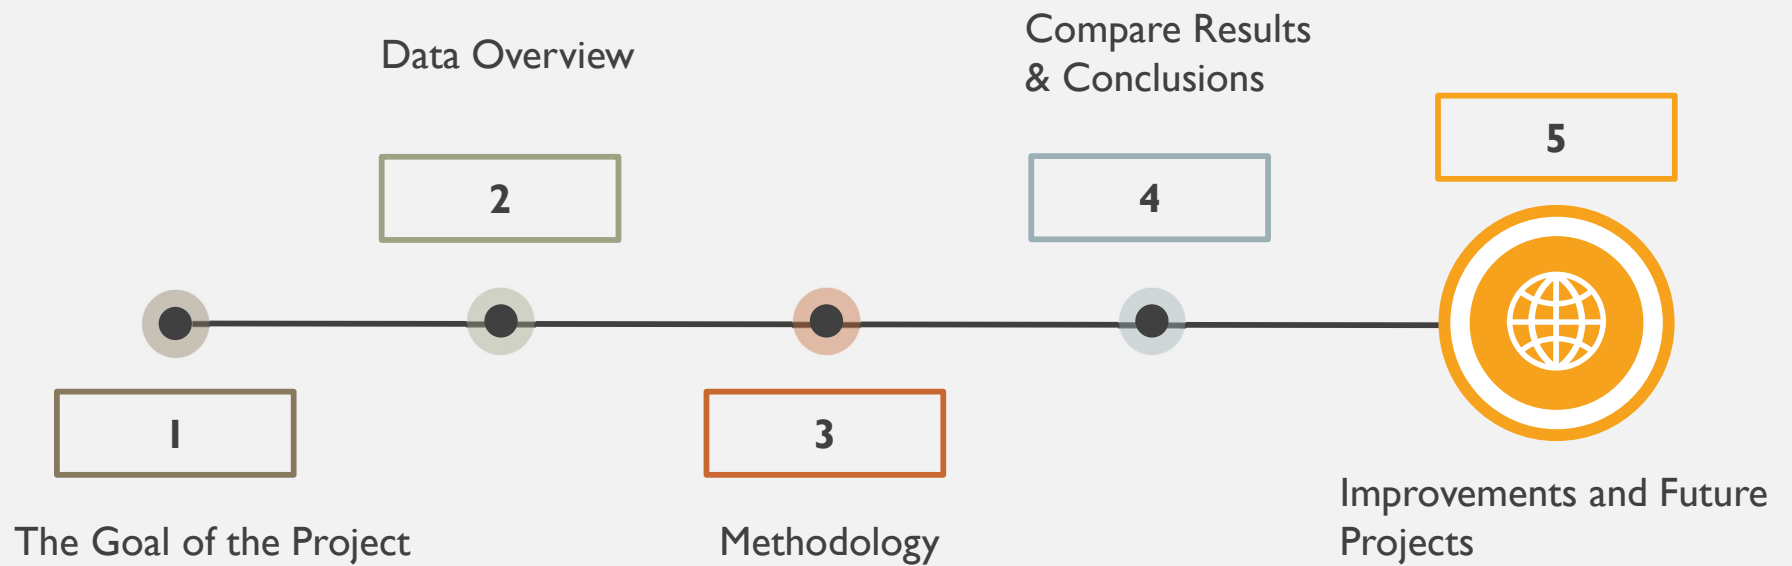


BANKING: PREDICTING CUSTOMER SUBSCRIPTION TO TERM DEPOSITS

AGENDA LAYOUT



BUSINESS GOAL



Problem

Banks need to avoid spending too much on individuals that will not subscribe to a term deposit.



Business Goal

To predict whether if a customer subscribes to a term deposit or not by using previous marketing campaign data.



Stakeholders

Customers, Marketing Team, Service Employees



Opportunities

Gain new customers and lower marketing cost by having more efficient marketing results



Challenge

Lose potential customers if the model predicts incorrectly

TERM DEPOSIT

What is it?

- Time deposit or fixed deposit or CD
- Usually with a financial institution
- Specific maturity date (commonly referred to as its “term”)
- Higher interest rate
- Cannot be withdrawn anytime (penalty)

Why does it matter to the bank?

- Knowing that term deposits allow banks to hold onto a deposit for a specific amount of time, so banks can invest in higher gain financial products to make a profit.
- In addition, banks also hold better chance to persuade term deposit clients into buying other products such as funds or insurance to further increase their revenues. The bank would like to identify existing clients that have higher chance to subscribe for a term deposit and focus marketing efforts on such clients.

DATA OVERVIEW



Source:

Direct telemarketing
campaign data from a
financial institution

of rows: 45,212
of Attributes: 17
with one decision
attribute

Input variables:

ID, Age, Job, Marital
Status, Education,
Default, Balance,
Housing, Loan, Contact,
Day, Month, Duration,
Campaign, pdays,
previous, poutcome.

Output variable: y
(subscribed/Not
subscribed)

Data Cleaning:

Eliminate null values,
ignore data points that
do not make
sense (age > 150 and age
< 0)

Normalize data and
change data type into
format usable by
analysis tool

EXPLORATORY DATA ANALYSIS



Mean age is 41 (min is 18 and max is 95)





Mean balance is \$1362 and standard deviation is high, So customer has a varying level of account balance.



The highest correlation coefficient is 0.39, which is between y (subscribed/not subscribed) and duration.

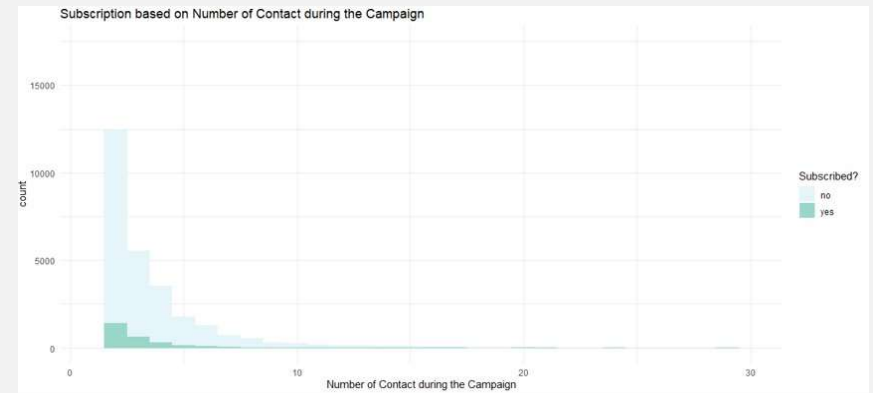
There is no multicollinearity as none of the dependent variables are highly correlated with each other



DATA VISUALIZATION

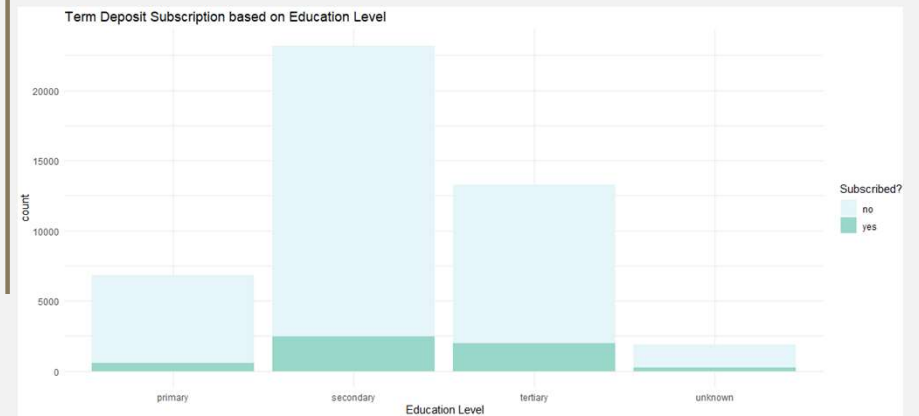
Subscription based on Number of Contact during Campaign

People that were going to subscribe did not receive many campaigns



Term Deposit Subscription based on Educational Level

People with higher education were more likely to subscribe to a term deposit



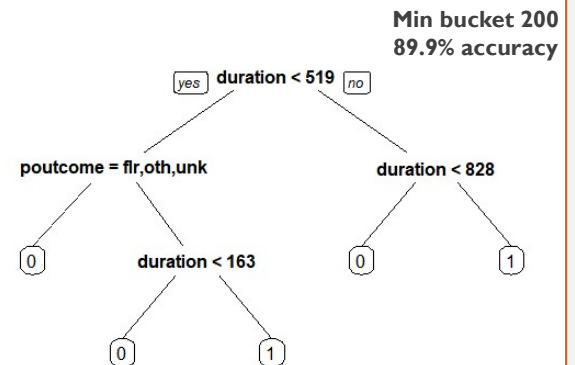
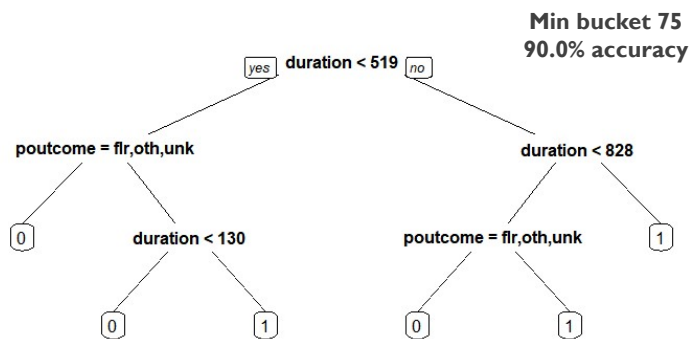
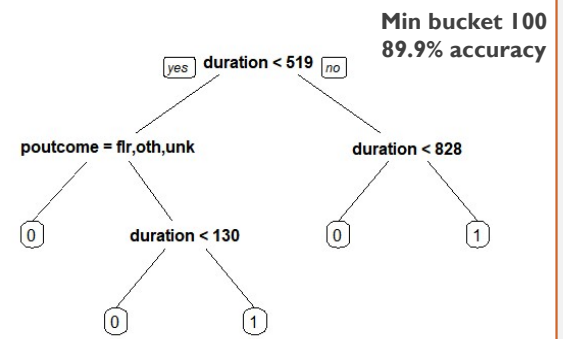
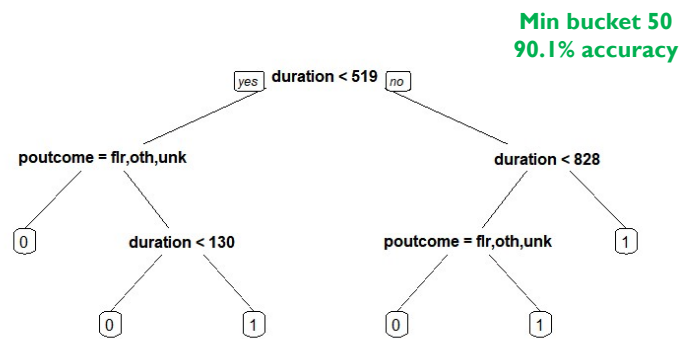
CART ANALYSIS

Different Trees

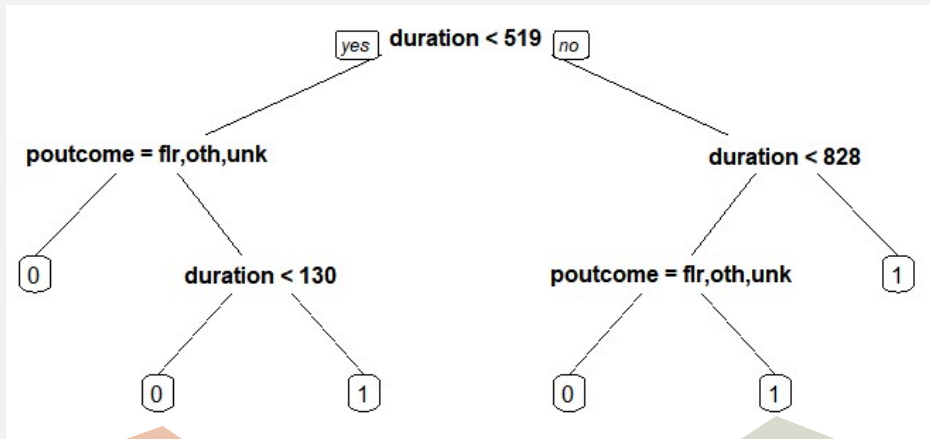
Changing the minimum size bucket alters the analysis, so we tried a few different groupings

Min Bucket=50

Ultimately, we agreed that a minimum grouping of 50 people was the most accurate



CART ANALYSIS




For customers, whose last contact was under 519 seconds, the outcome of the marketing campaign was not a success, and lastly, the last contact duration was under 130 seconds, a term deposit is unlikely (<50%)

For customers, whose last contact was over 519 seconds (and over 828 seconds) and the outcome of the marketing campaign was a success, a term deposit is likely (>50%)

Accuracy: 0.90

Sensitivity: 0.37



	Pred_0	Pred_1
Actual_0	11627	333
Actual_1	1012	587

LOGISTIC REGRESSION ANALYSIS


Deviance Residuals:

Min	1Q	Median	3Q	Max
-5.6225	-0.4574	-0.2937	-0.1613	3.6717

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-2.302e+00	9.236e-02	-24.922	< 2e-16 ***
age	-1.208e-03	1.736e-03	-0.695	0.487
default	-1.239e-01	1.781e-01	-0.695	0.487
balance	2.451e-05	5.065e-06	4.839	1.31e-06 ***
housing	-1.006e+00	4.274e-02	-23.549	< 2e-16 ***
loan	-7.454e-01	6.594e-02	-11.304	< 2e-16 ***
contacttelephone	-6.111e-02	7.735e-02	-0.790	0.430
contactunknown	-1.337e+00	6.575e-02	-20.337	< 2e-16 ***
day	-3.302e-03	2.371e-03	-1.393	0.164
duration	3.917e-03	6.983e-05	56.087	< 2e-16 ***
campaign	-1.420e-01	1.170e-02	-12.132	< 2e-16 ***
pdays	2.173e-03	1.926e-04	11.279	< 2e-16 ***
previous	6.211e-02	8.741e-03	7.106	1.20e-12 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1



	Pred_0	Pred_1
Actual_0	9719	259
Actual_1	959	363

Accuracy: 0.89

Sensitivity: 0.36

RANDOM FOREST ANALYSIS

Model Output

Confusion Matrix and Statistics

	Reference	
Prediction	0	1
0	9731	798
1	247	524

Accuracy : 0.9075
95% CI : (0.902, 0.9128)
No Information Rate : 0.883
P-Value [Acc > NIR] : < 2.2e-16

Kappa : 0.4536


Mcnemar's Test P-Value : < 2.2e-16

Sensitivity : 0.39637
Specificity : 0.97525
Pos Pred Value : 0.67964
Neg Pred Value : 0.92421
Prevalence : 0.11699
Detection Rate : 0.04637
Detection Prevalence : 0.06823
Balanced Accuracy : 0.68581

'Positive' Class : 1


Accuracy: 0.9075

Variable Importance



Sensitivity: 0.40

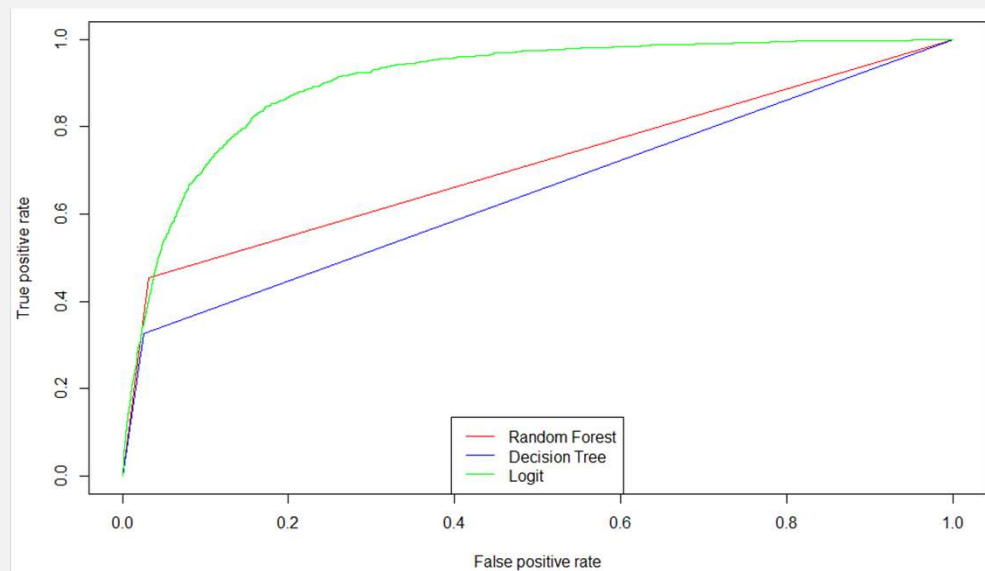
ROC Curve



	Pred_0	Pred_1
Actual_0	9731	247
Actual_1	798	524

COMPARISON RESULTS

Model	Sensitivity	AUC	Accuracy	Specificity
Random Forest	0.40	0.69	0.91	0.98
Logistic Regression	0.36	0.87	0.89	0.97
CART	0.37	0.65	0.90	0.97



Since our purpose is to generate a model for bank marketing purposes, a false-negative is way more harmful to the marketing strategy than a false positive. Since the goal of marketing is to get as many positives as possible, we are going to use sensitivity as the metric to compare our models' performance.

Based on sensitivity and accuracy, the Random Forrest model proves the most useful in identifying customers interested in getting a term deposit

POTENTIAL IMPROVEMENTS




Data Imbalance: (88.3% do not subscribe, 11.7% subscribe). An imbalanced dataset may lead to inaccurate predictions

Fix: Random Sampling Technique by sampling data from minority class and duplicate it to create more samples



Problem: What can we do to have a more efficient/flexible model?

Fix: Improve our models by using gradient boosting framework
-> XGBoost, SVM



Problem: Too many variables that could cost a lot to analyze the dataset.

Fix: Use Feature Selection & Grid Search to reduce the number of input variables (ID, day, month)

FUTURE PROJECTS



Seasonality

What is the best season (month/day) to contact the potential client?



Scoring Prospects

Use scoring/rating to measure a prospect's intent and interest. Their scores based on how they interact with the bank and change over time



Customer Segmentation

Segment customers into small groups to develop better marketing strategies for each group





A/B Testing

Run two versions of marketing campaigns to collect data to see which one performs better

THANK YOU

APPENDIX



CART CODE

```
#Loading Libraries
```

```
library(caTools)
```

```
library(rpart)
```

```
library(rpart.plot)
```

```
#Data Cleaning (removing weird ages and NAs)
```

```
FinalData <- FinalData[-c(which(FinalData$age <= 0 | FinalData$age > 150)),]
```

```
FinalData <- na.omit(FinalData)
```

```
FinalData$y[which(FinalData$y == "yes")] <- 1
```

```
FinalData$y[which(FinalData$y == "no")] <- 0
```

```
FinalData$y <- as.numeric(FinalData$y)
```

```
#Training and Testing Dataset
```

```
set.seed(88)
```

```
split <- sample.split(FinalData$Id, SplitRatio = 0.7)
```

```
Train <- subset(FinalData, split==TRUE)
```

```
Test <- subset(FinalData, split==FALSE)
```

```
#Set baseline
```

```
nrow(Train)
```

```
sum(Train$y)
```

```
sum(Train$y)/nrow(Train)
```

```
#CART Analysis
```

```
Tree_min50= rpart(y~ age + job + marital + education + default + balance + housing + loan + contact + day + month + duration + campaign + pdays + previous + poutcome, method = "class", data=Train, minbucket=50)
```

```
prp(Tree_min50)
```

```
Tree_min50Predict = predict(Tree_min50, newdata = Test, type="class")
```

```
tbl_min50 = table(Test$y, Tree_min50Predict)
```

```
tbl_min50
```

```
sum(diag(tbl_min50))/sum(tbl_min50)
```

```
Tree_min100= rpart(y~ age + job + marital + education + default + balance + housing + loan + contact + day + month + duration + campaign + pdays + previous + poutcome, method = "class", data=Train, minbucket=100)
```

```
prp(Tree_min100)
```

```
Tree_min100Predict = predict(Tree_min100, newdata = Test, type="class")
```

```
tbl_min100 = table(Test$y, Tree_min100Predict)
```

```
tbl_min100
```

```
sum(diag(tbl_min100))/sum(tbl_min100)
```

```
Tree_min200= rpart(y~ age + job + marital + education + default + balance + housing + loan + contact + day + month + duration + campaign + pdays + previous + poutcome, method = "class", data=Train, minbucket=200)
```

```
prp(Tree_min200)
```

```
Tree_min200Predict = predict(Tree_min200, newdata = Test, type="class")
```

```
tbl_min200 = table(Test$y, Tree_min200Predict)
```

```
tbl_min200
```

```
sum(diag(tbl_min200))/sum(tbl_min200)
```

```
Tree_min75= rpart(y~ age + job + marital + education + default + balance + housing + loan + contact + day + month + duration + campaign + pdays + previous + poutcome, method = "class", data=Train, minbucket=75)
```

```
prp(Tree_min75)
```

```
Tree_min75Predict = predict(Tree_min75, newdata = Test, type="class")
```

```
tbl_min75 = table(Test$y, Tree_min75Predict)
```

```
tbl_min75
```

```
sum(diag(tbl_min75))/sum(tbl_min75)
```

LOGIT CODE

```
data = read.csv("Assignment-2_Data.csv")
summary(data)
str(data)

# remove ages that do not make sense
data <- data[-c(which(data$age <= 0 | data$age>
150)),]

# remove na values
bank <- na.omit(data)
bank$y[bank$y == "no"] = 0
bank$y[bank$y == "yes"] = 1
bank$housing[bank$housing == "no"] = 0
bank$housing[bank$housing == "yes"] = 1
bank$loan[bank$loan == "no"] = 0
bank$loan[bank$loan == "yes"] = 1
bank$default[bank$default == "no"] = 0
bank$default[bank$default == "yes"] = 1

bank$housing <- as.integer(bank$housing)
bank$default <- as.integer(bank$default)
bank$loan <- as.integer(bank$loan)
bank$y <- as.integer(bank$y)
```

```
# Creating Training and Testing Sets
library(caTools)
# Normalization
library(caret)
str(bank)
#mean and sd of each variable
preproc = preProcess(bank)
#normalize the data
BankNorm = predict(preproc, bank)
# Creating Training and Testing Sets
set.seed(88)
split = sample.split(BankNorm$y, SplitRatio = 0.7
5)
Train = subset(bank, split==TRUE)
Test = subset(bank, split==FALSE)

# Building a Logistic Regression Model
str(Train)
bankLog = glm(y ~ age+default+balance+housing
+loan+contact+day+duration+campaign+pdays+p
revious, data = Train, family=binomial)
summary(bankLog)
```

```
# Evaluating the Model

predicted_values<-ifelse(predict(bankLog,type="r
esponse", newdata = Test)>0.4,1,0)
actual_values<-Test$y
conf_matrix<-table(predicted_values,actual_valu
es)
conf_matrix
specificity(conf_matrix)
sensitivity(conf_matrix)
library(ROCR)
PredictTest = predict(bankLog, type="response",
newdata = Test)
summary(PredictTest)
tbl = table(Test$y, PredictTest > 0.4)
sum(diag(tbl))/sum(tbl)
ROCRpred = prediction(PredictTest, Test$y)
ROCCurve = performance(ROCRpred, "tpr", "fp
r")
plot(ROCCurve)
plot(ROCCurve, colorize=TRUE, print.cutoffs.at
=seq(0,1,0.1), text.adj=c(-0.2,0.7))
as.numeric(performance(ROCRpred, "auc")@y.v
alues) # AUC value
```

ROC CODE

```
library(pacman)
library(janitor)
# Reading in the data
term <- clean_names(read.csv("Assignment-2_Data.csv"))
term <- term[-c(which(term$age <= 0 | term$age>150)),]
term <- na.omit(term)
str(term)
summary(term)
term$y <- ifelse(term$y=="yes", 1,0)
term$y <- as.integer(term$y)
term$housing <- ifelse(term$housing == "yes", 1,0)
term$housing <- as.integer(term$housing)
term$default <- ifelse(term$default=="yes", 1,0)
term$default <- as.integer(term$default)
term$loan <- ifelse(term$loan=="yes", 1,0)
term$loan <- as.integer(term$loan)
term[apply(term, is.character)] <- lapply(term[apply(term, is.character)],
                                           as.factor)
# calling the library need for splitting
library(caTools)
set.seed(88)
# splitting into training and testing
split = sample.split(term, SplitRatio = 0.75)
train = subset(term, split == TRUE)
test = subset(term, split == FALSE)
# Random Forest
library(randomForest)
train$y <- as.factor(train$y)
test$y <- as.factor(test$y)
TermForrest <- randomForest(y ~ age + job + marital + education + default + balance + housing +
                             loan + contact + day + duration + campaign + pdays + previous
                             + month + poutcome, data=train, ntree=200, nodesize=15)
```

```
TermPredict <- predict(TermForrest, newdata=test)
tbl <- table(test$y, TermPredict)
sum(diag(tbl))/sum(tbl)

# Building a CART model
library(rpart)
library(rpart.plot)
TermTree = rpart(y ~ age + job + marital + education + default
                  + balance + housing +
                  loan + contact + day + duration + campaign + pday
                  s + previous
                  + month + poutcome, method="class", data=train,
                  minbucket=100)
prp(TermTree)

# automatically assumes a threshold of .5
TermPredict2 = predict(TermTree, newdata=test, type="class")
tbl2 <- table(test$y, TermPredict)
sum(diag(tbl))/sum(tbl)
```

```
# Running a logistic Regression
bankLog = glm(y ~ age+marital+education+contact+month+poutcome+default+balance+housing+loan+day+duration+campaign
+pdays+previous , data = train, family=binomial)
summary(bankLog)
```

```
# Evaluating the Model
PredictTrain = predict(bankLog, type="response")
table <- table(train$y, PredictTrain > 0.5)
sum(diag(table))/sum(table)
```

```
# Testing model on new data
PredictTest = predict(bankLog, type="response", newdata=test)
table2 <- table(test$y, PredictTest > 0.5)
sum(diag(table2))/sum(table2)
# ROC Curve
library(ROCR)
ROCRpredrf = as.numeric(TermPredict)
ROCRpred = prediction(ROCRpredrf, test$y)
ROCCurve = performance(ROCRpred, "tpr", "fpr")
as.numeric(performance(ROCRpred, "auc")@y.values)
# ROC Curve 2
ROCRpredrf2 = as.numeric(TermPredict2)
ROCRpred2 = prediction(ROCRpredrf2, test$y)
ROCCurve2 = performance(ROCRpred2, "tpr", "fpr")
as.numeric(performance(ROCRpred2, "auc")@y.values)
# ROC Curve 3
ROCRpred3 = prediction(PredictTest, test$y)
ROCCurve3 = performance(ROCRpred3, "tpr", "fpr")
as.numeric(performance(ROCRpred3, "auc")@y.values)
plot(ROCCurve, col = 'red')
plot(ROCCurve2, add = TRUE, col = 'blue')
plot(ROCCurve3, add = TRUE, col = 'green')
legend("bottom", c("Random Forest","Decision Tree","Logit"), lty=1,
      col = c("red", "blue", "green"))
```

RANDOM FOREST CODE

```
data = read.csv("Assignment-2_Data.csv")
summary(data)
# remove ages that do not make sense
data <- data[-c(which(data$age <= 0 | data$age>150)),
]
```

```
# remove na values
bank <- na.omit(data)
bank$y[bank$y == "no"] = 0
bank$y[bank$y == "yes"] = 1
bank$housing[bank$housing == "no"] = 0
bank$housing[bank$housing == "yes"] = 1
bank$loan[bank$loan == "no"] = 0
bank$loan[bank$loan == "yes"] = 1
bank$default[bank$default == "no"] = 0
bank$default[bank$default == "yes"] = 1
```

```
bank$housing <- as.integer(bank$housing)
bank$default <- as.integer(bank$default)
bank$loan <- as.integer(bank$loan)
bank$y <- as.integer(bank$y)
bank[sapply(term, is.character)] <- lapply(bank[sapply(bank, is.character)],
                                             as.factor)
```

```
# Creating Training and Testing Sets
library(caTools)
```

```
# Normalization
library(caret)
```

```
#mean and sd of each variable
preproc = preProcess(bank)
```

```
#normalize the data
BankNorm = predict(preproc, bank)
```

```
# Creating Training and Testing Sets
set.seed(88)
split = sample.split(BankNorm$y, SplitRatio = 0.75)
Train = subset(bank, split==TRUE)
Test = subset(bank, split==FALSE)
```

```
#Random Forest Model
library(randomForest)
```

```
Train$y = as.factor(Train$y)
Test$y = as.factor(Test$y)
```

```
#Random Forest Model
rfmdel <- randomForest(y ~age+default+balance+housing+loan+contact+day+duration+campaign+pdays+previous+marital+education+poutcome+job+month,data = Train)
varImpPlot(rfmdel) #To Check Variable Importance
prediction_rf <- predict(rfmdel,Test)
prediction_rf
```

```
#Confusion matrix to validate it
conf_mat <- confusionMatrix(prediction_rf,Test$y)
```

```
#Table for Validation
rftable = table(Test$y,prediction_rf)
Accuracy = sum(diag(rftable))/sum(rftable)
Accuracy
```

```
#ROC Curve
library(ROCR)
ROCRpredrf = as.numeric(predict(rfmdel,Test,type="response"))
ROCRf = prediction(ROCRpredrf,Test$y)
ROCCurvef = performance(ROCRf,measure = "tpr", x.measure = "fpr")
plot(ROCCurvef)
plot(ROCCurvef, colorize=TRUE, print.cutoffs.at=seq(0,1,0.1), text.adj=c(-0.2,0.7))
AUC_rf <- as.numeric(performance(ROCRf, "auc")@y.values) # AUC value
AUC_rf
```